



# Lets talk about the filesystem

---

- Recall the boot process...
  - BIOS starts POST
  - Finds DRIVES using BIOS sequence
  - Finds BOOTABLE MEDIA
    - Based on MBR (512 Bytes)



# MBR structure

---

- ❑ MBR executable code starts at offset 0x0000, total 446 bytes
  - The MBR messages start at offset 0x008b
- ❑ The partition table starts at offset 0x01be
  - 64 bytes in four 16 bytes sections
  - 1<sup>st</sup> partition Entry starts at offset 0x01be
  - 2<sup>nd</sup> partition entry starts at offset 0x01ce
  - 3<sup>rd</sup> partition entry starts at offset 0x01de
  - 4<sup>th</sup> partition entry starts at offset 0x01ee
- ❑ The signature is at offset 0x01fe, 2 bytes (55AAh)
- ❑ Total 512 bytes



# Disk Structure

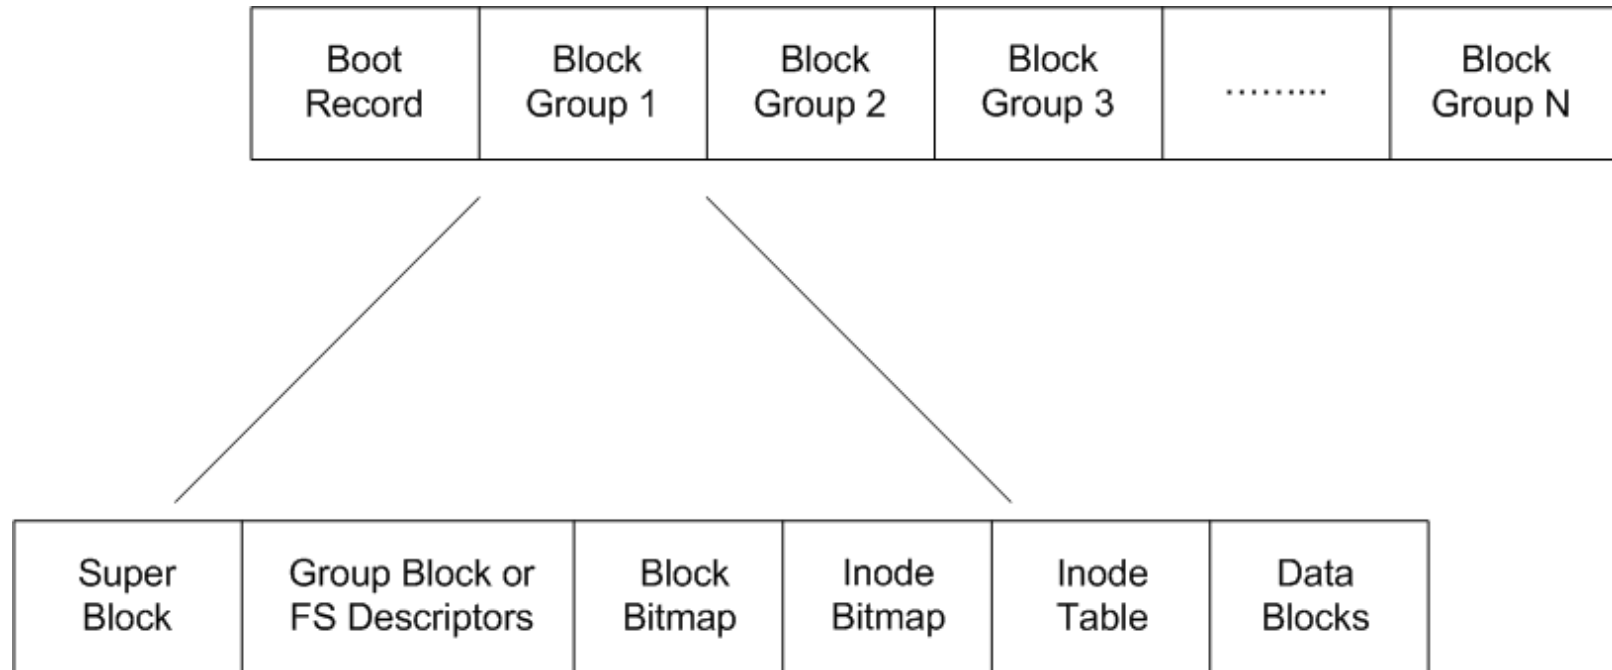
---

- ❑ Disks are partitioned
- ❑ Each partition contains the following areas:
  - An optional boot block
  - A super block defining the boundaries of the other areas.
  - Block Bitmap
  - Inode Bitmap
  - A set of file information blocks known as I-nodes
  - The data blocks (free and used intermingled)

# Functionality

---

## □ ext2/ext3 filesystem structure



[illegible]

### I-Node Bitmap



# superblock

---

- ❑ Superblock contains
  - Magic Number
    - ❑ For EXT2, it is 0xEF53
  - Mount Count and Maximum Mount Count
  - Block Size, for example 4096B
  - Inode count and Block count
  - Number of free disk blocks
  - Number of free inodes on the system
  - First inode
    - ❑ This is the inode number of the first inode in the file system. The first inode in an EXT2 root file system would be the directory entry for the ‘/’ directory



# Inodes

---

- Contain META data (info about files)
  - type
  - Owner
  - Permission bits
  - MAC times
  - Links to the file (link count)
  - Data block addresses





# I-nodes System V

---

- An I-node consists of 16 32 bit words
  - Word 0 is split into two 16 bit words
    - These contain the file type and mode and the link count (more later).
  - Word 1 is split into two 16 bit words
    - The user and group ids
  - Word 2 gives the file size in bytes
    - Implying a maximum filesize of 4 GB



## I-nodes System V <sub>cont.</sub>

---

- Word 13 contains the time of last access
- Word 14 contains the time of last data modification
- Word 15 contains the time of the last file status change.

Type/Mode	Link Count
User Id	Group Id
File size (in bytes)	
10 Data Block Addresses	
	Unused
Time last accessed	
Time last modified	
Time created	

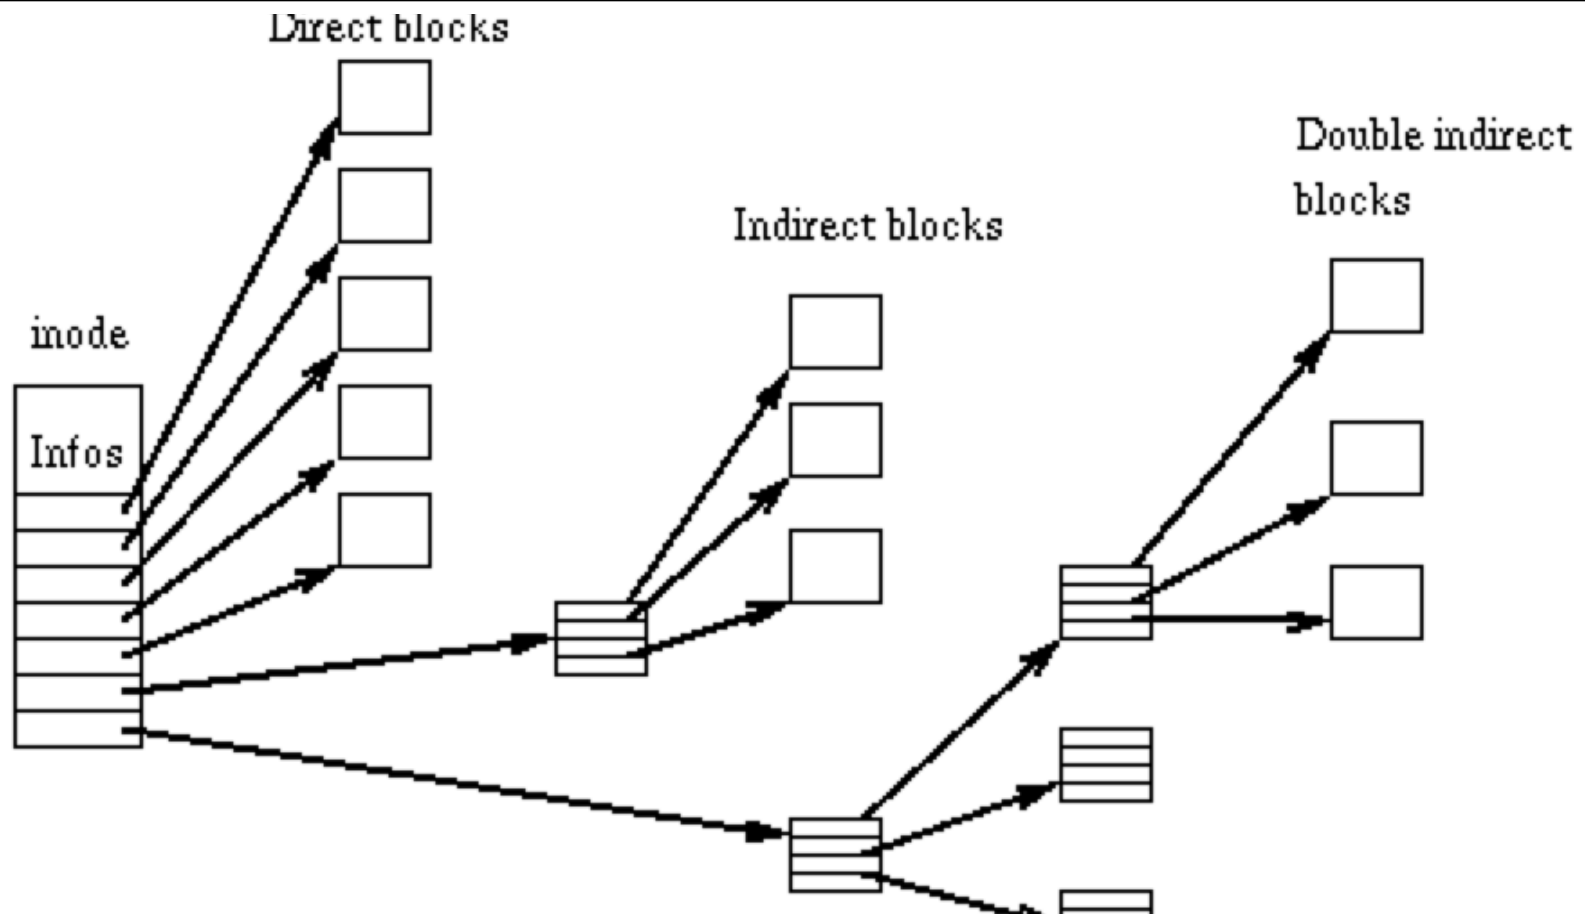
First Level Index Block Address

Second Level Index Block Address

Third Level Index Block Address

32 bits

# Inode pointer structure (From wikipedia)





# Directories

---

- Everything in Linux is a file
- Directory is a simple file whose data is a sequence of file entries which contains:
  - Inode number
  - Byte offset in directory (or the length of this entry)
  - File name

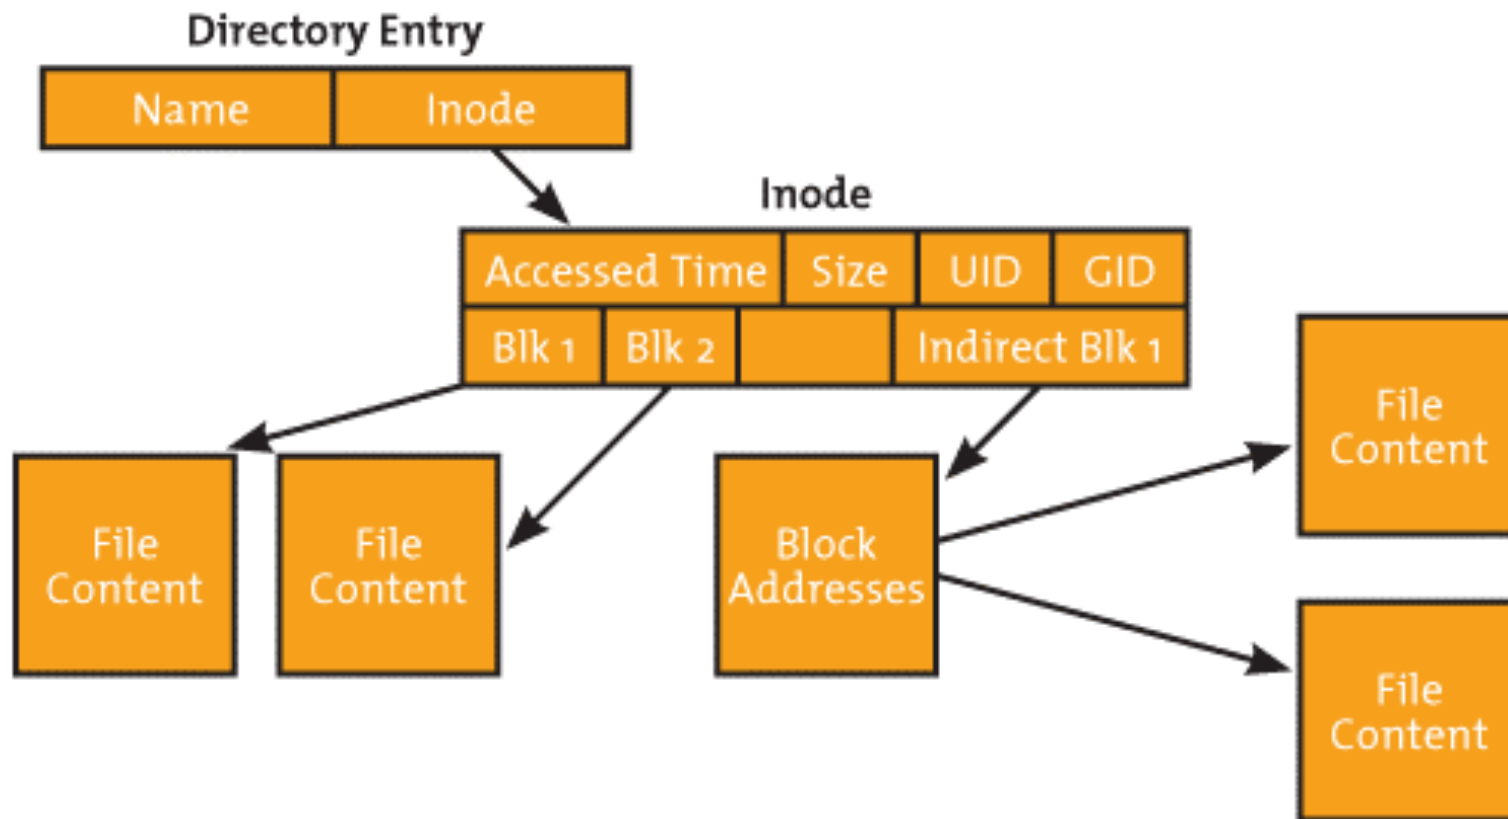


# Directory example

---

Byte offset in directory	Inode number	File name
0	70	.
16	35	..
32	123	file1
48	345	file2
64	90	file3

# Directory entry



Carrier B., Why Recovering a Deleted Ext3 File  
Is Difficult . . .